

Automatic and Interpretable Machine Learning with H2O & LIME

Code Breakfast at GoDataDriven

Jo-fai (Joe) Chow - joe@h2o.ai - @matlabulous

Download: https://github.com/woobe/lime_water/ or
bit.ly/joe_lime_water

About Joe

- My name is Jo-fai
- Majority of my British friends cannot remember Jo-fai
- Joe is the solution
- Data Scientist at H2O.ai
- For a very long time, I was the only H2O person in UK ...
- Community Manager / Sales Engineer / Photographer / SWAG Distributor



Agenda

- **Introduction**
 - Why?
 - Interpretable Machine Learning
 - LIME Framework
 - Automatic Machine Learning
 - H2O AutoML
- **Worked Examples**
 - Regression
 - Classification
- **Other Stuff + Q & A**



+

H₂O.ai

Acknowledgement

- **Marco Tulio Ribeiro:** Original LIME Framework and Python package
- **Thomas Lin Pedersen:** LIME R package
- **Matt Dancho:** LIME + H2O AutoML example + LIME R package improvement
- **Kasia Kulma:** LIME + H2O AutoML example
- My H2O colleagues **Erin LeDell, Ray Peck, Navdeep Gill** and many others for AutoML

Why?

Why Should I Trust Your Model?



System that performs behaviour but you don't know how it works

Interpretable Machine Learning




Figure 1. Explaining individual predictions to a human decision-maker. Source: Marco Tulio Ribeiro.




Figure 2. Explaining a model to a human decision-maker. Source: Marco Tulio Ribeiro.

The LIME Framework

Local Interpretable Model-agnostic Explanations

A framework for interpretability


Complexity of learned functions:

- Linear, monotonic
- Nonlinear, monotonic
- Nonlinear, non-monotonic



Scope of interpretability:

Global vs. local



Enhancing trust and understanding: the mechanisms and results of an interpretable model should be both transparent AND dependable.



Application domain:

Model-agnostic vs. model-specific




Linear Models

Exact explanations for **approximate** models.



Machine Learning

Approximate explanations for **exact** models.



How does LIME work?

Theory


- LIME approximates model locally as logistic or linear model
- Repeats process many times
- Output features that are most important to local models

Outcome

- Approximate reasoning
- Complex models can be interpreted
 - Neural nets, Random Forest, Ensembles etc.

Automatic Machine Learning

Typical Enterprise Machine Learning Workflow



H2O AutoML

H2O's AutoML can be used for automating a large part of the machine learning workflow, which includes automatic training and tuning of many models within a user-specified time-limit. The user can also use a performance metric-based stopping criterion for the AutoML process rather than a specific time constraint. Stacked Ensembles will be automatically trained on the collection individual models to produce a highly predictive ensemble model which, in most cases, will be the top performing model in the AutoML Leaderboard.

R Interface

```
aml = h2o.automl(x = x, y = y,  
                  training_frame = train,  
                  max_runtime_secs = 3600)
```

Python Interface

```
aml = H2OAutoML(max_runtime_secs = 3600)  
aml.train(x = x, y = y,  
          training_frame = train)
```

Web Interface

H2O Flow

Lime Water

Lime Water in R

LIME

```
# Install 'lime' from CRAN  
install.packages('lime')
```

or

```
# Install development version from GitHub  
devtools::install_github('thomasp85/lime')
```

H2O

```
# Install 'h2o' from CRAN  
install.packages('h2o')
```

or

```
# Install latest stable release from H2O's  
# website www.h2o.ai/download/  
# Latest Version = 3.18.0.1  
# (as of 19-Feb-2018)  
install.packages("h2o", type="source",  
repos="http://h2o-release.s3.amazonaws.com  
/h2o/rel-wolpert/1/R")
```

Regression Example

Regression Example: Boston Housing

Data Set Characteristics:

- Number of Instances: 506
- Number of Attributes: 13 numeric/categorical predictive
- Median Value (attribute 14) is the target
- Attribute Information (in order):
 - CRIM per capita crime rate by town
 - ZN proportion of residential land zoned for lots over 25,000 sq.ft.
 - INDUS proportion of non-retail business acres per town
 - CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
 - NOX nitric oxides concentration (parts per 10 million)
 - RM average number of rooms per dwelling
 - AGE proportion of owner-occupied units built prior to 1940
 - DIS weighted distances to five Boston employment centres
 - RAD index of accessibility to radial highways
 - TAX full-value property-tax rate per \$10,000
 - PTRATIO pupil-teacher ratio by town
 - B $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town
 - LSTAT % lower status of the population
 - MEDV Median value of owner-occupied homes in \$1000's
- Creator: Harrison, D. and Rubinfeld, D.L.
- Source: <http://archive.ics.uci.edu/ml/datasets/Housing>

Regression Example: Boston Housing

```
library(mlbench) # for dataset  
data("BostonHousing")  
dim(BostonHousing)
```

```
## [1] 506 14
```

```
# First six samples  
knitr::kable(head(BostonHousing), format = "html")
```

crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv
0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2
0.02985	0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7	394.12	5.21	28.7

Boston Housing (Simple Split)

```
# Define features
features = setdiff(colnames(BostonHousing), "medv")
features

## [1] "crim"      "zn"        "indus"     "chas"      "nox"       "rm"        "age"
## [8] "dis"        "rad"       "tax"        "ptratio"   "b"         "lstat"

# Pick four random samples for test dataset
set.seed(1234)
row_test_samp = sample(1:nrow(BostonHousing), 4)

# Train
x_train = BostonHousing[-row_test_samp, features]
y_train = BostonHousing[-row_test_samp, "medv"]

# Test
x_test = BostonHousing[row_test_samp, features]
y_test = BostonHousing[row_test_samp, "medv"]
```

Build a Random Forest (RF)

```
library(caret) # ML framework  
library(doParallel) # parallelisation
```

```
# Train a Random Forest using caret  
cl = makePSOCKcluster(8)  
registerDoParallel(cl)  
set.seed(1234)  
model_rf =  
  caret::train(  
    x = x_train,  
    y = y_train,  
    method = "rf",  
    tuneLength = 3,  
    trControl = trainControl(method = "cv")  
  )  
stopCluster(cl)
```

```
# Print model summary  
model_rf
```

```
## Random Forest  
##  
## 502 samples  
## 13 predictor  
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold)  
## Summary of sample sizes: 453, 451, 453, 451, 452, 452, ...  
## Resampling results across tuning parameters:  
##  
##     mtry   RMSE    Rsquared    MAE  
##     2      3.532344  0.8715278  2.361838  
##     7      3.204915  0.8874722  2.185825  
##    13     3.256840  0.8798610  2.230762  
##  
## RMSE was used to select the optimal model using the small  
## The final value used for the model was mtry = 7.
```

RF: Making Prediction

```
# Using the Random Forest model to make predictions on test set  
yhat_test = predict(model_rf, x_test)
```

```
# Create a new data frame to compare target (medv) and predictions  
d_test = data.frame(x_test,  
                     medv = y_test,  
                     predict = yhat_test,  
                     row.names = NULL)  
knitr::kable(d_test, format = "html")
```

crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv	predict
0.01432	100	1.32	0	0.411	6.816	40.5	8.3248	5	256	15.1	392.90	3.95	31.6	31.63876
0.36920	0	9.90	0	0.544	6.567	87.3	3.6023	4	304	18.4	395.69	9.28	23.8	24.09371
0.04932	33	2.18	0	0.472	6.849	70.3	3.1827	7	222	18.4	396.90	7.53	28.2	29.78971
0.26938	0	9.90	0	0.544	6.266	82.8	3.2628	4	304	18.4	393.39	7.90	21.6	22.69258

RF: LIME Steps 1 and 2

```
# Step 1: Create an 'explainer' object using training data and model  
explainer = lime::lime(x = x_train, model = model_rf)
```

```
# Step 2: Turn 'explainer' into 'explanations' for test set  
explanations = lime::explain(x = x_test,  
                           explainer = explainer,  
                           n_permutations = 5000,  
                           feature_select = "auto",  
                           n_features = 5)
```


RF: LIME Explanations

```
head(explanations, 5) #LIME Pred: 36.59, Random Forest Pred: 31.64, R^2 = 0.65
```

```
##   model_type case  model_r2 model_intercept model_prediction feature
## 1 regression   58 0.6533429      20.43179     36.59478      rad
## 2 regression   58 0.6533429      20.43179     36.59478      rm
## 3 regression   58 0.6533429      20.43179     36.59478    lstat
## 4 regression   58 0.6533429      20.43179     36.59478      dis
## 5 regression   58 0.6533429      20.43179     36.59478      tax
##   feature_value feature_weight feature_desc
## 1      5.0000      0.1918038 4 < rad ≤ 5
## 2      6.8160      8.5174963 6.62 < rm
## 3      3.9500      7.6363579 lstat ≤ 6.95
## 4      8.3248     -1.0748818 5.19 < dis
## 5     256.0000      0.8922157 tax ≤ 279
##
## 1 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 2 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 3 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 4 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 5 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
##   prediction
## 1 31.63876
## 2 31.63876
## 3 31.63876
## 4 31.63876
## 5 31.63876
```

RF: LIME Visualisation

```
# Step 3: Visualise explanations  
lime::plot_features(explanations, ncol = 4)
```



H2O AutoML

```
# Start a local H2O cluster (JVM)
library(h2o)
h2o.init(nthreads = -1)

## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      3 days 9 hours
##   H2O cluster timezone:    Europe/London
##   H2O data parsing timezone: UTC
##   H2O cluster version:     3.18.0.1
##   H2O cluster version age: 6 days
##   H2O cluster name:        H2O_started_from_R_jofaichow_ydb410
##   H2O cluster total nodes: 1
##   H2O cluster total memory: 3.29 GB
##   H2O cluster total cores: 8
##   H2O cluster allowed cores: 8
##   H2O cluster healthy:     TRUE
##   H2O Connection ip:       localhost
##   H2O Connection port:     54321
##   H2O Connection proxy:    NA
##   H2O Internal Security:  FALSE
##   H2O API Extensions:    XGBoost, Algos, AutoML, Core V3, Core V4
##   R Version:               R version 3.4.3 (2017-11-30)
```

Prepare H2O Data Frames

```
# Prepare Data
h_train = as.h2o(BostonHousing[-row_test_samp,])
h_test = as.h2o(BostonHousing[row_test_samp,])

head(h_test)

##      crim   zn indus chas   nox     rm    age     dis   rad tax ptratio      b
## 1 0.01432 100  1.32     0 0.411 6.816 40.5 8.3248     5 256  15.1 392.90
## 2 0.36920    0  9.90     0 0.544 6.567 87.3 3.6023     4 304  18.4 395.69
## 3 0.04932   33  2.18     0 0.472 6.849 70.3 3.1827     7 222  18.4 396.90
## 4 0.26938    0  9.90     0 0.544 6.266 82.8 3.2628     4 304  18.4 393.39
##      lstat medv
## 1 3.95 31.6
## 2 9.28 23.8
## 3 7.53 28.2
## 4 7.90 21.6
```

Train Multiple H2O Models

```
# Train multiple H2O models with a simple API
# Stacked Ensembles will be created from those H2O models
# You tell H2O 1) how much time you have and/or 2) how many models do you want
model_automl = h2o.automl(x = features,
                           y = "medv",
                           training_frame = h_train,
                           nfolds = 5,
                           max_runtime_secs = 120, # time
                           max_models = 20,       # max models
                           stopping_metric = "RMSE",
                           seed = 1234)
```

H2O: AutoML Model Leaderboard

```
# Print out leaderboard
model_automl@leaderboard

##                                     model_id
## 1 StackedEnsemble_BestOfFamily_0_AutoML_20180219_195835
## 2             GBM_grid_0_AutoML_20180219_195835_model_0
## 3 StackedEnsemble_AllModels_0_AutoML_20180219_195835
## 4             GBM_grid_0_AutoML_20180219_195835_model_1
## 5             GBM_grid_0_AutoML_20180219_195835_model_3
## 6                 DRF_0_AutoML_20180219_195835
##   mean_residual_deviance      rmse       mae     rmsle
## 1          10.84287 3.292852 2.151855 0.140915
## 2          10.86044 3.295518 2.224282 0.145063
## 3          10.91543 3.303851 2.163235 0.141070
## 4          11.88445 3.447383 2.285338 0.145858
## 5          12.12041 3.481438 2.324986 0.148829
## 6          12.22679 3.496683 2.339066 0.148301
##
## [22 rows x 5 columns]
```

H2O: Model Leader

```
# Best Model (either an individual model or a stacked ensemble)
model_automl@leader

## Model Details:
## =====
##
## H2OResponse: numeric
## Model ID: StackedEnsemble_BestOfFamily_0_AutoML_20180219_195835
## NULL
##
##
## H2OResponseMetrics: numeric
## ** Reported on training data. **
##
## MSE: 0.8388527
## RMSE: 0.915889
## MAE: 0.6740673
## RMSLE: 0.0451058
## Mean Residual Deviance : 0.8388527
##
##
## H2OResponseMetrics: numeric
## ** Reported on validation data. **
##
## MSE: 6.737158
## RMSE: 2.595604
## MAE: 1.74255
```

H2O: Making Prediction

```
# Using the best model to make predictions on test set  
yhat_test = h2o.predict(model_automl@leader, h_test)
```

```
# Create a new data frame to compare target (medv) and predictions  
d_test = data.frame(x_test,  
                     medv = y_test,  
                     predict = as.data.frame(yhat_test),  
                     row.names = NULL)  
knitr::kable(d_test, format = "html")
```

crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	lstat	medv	predict
0.01432	100	1.32	0	0.411	6.816	40.5	8.3248	5	256	15.1	392.90	3.95	31.6	31.02489
0.36920	0	9.90	0	0.544	6.567	87.3	3.6023	4	304	18.4	395.69	9.28	23.8	22.91670
0.04932	33	2.18	0	0.472	6.849	70.3	3.1827	7	222	18.4	396.90	7.53	28.2	30.13234
0.26938	0	9.90	0	0.544	6.266	82.8	3.2628	4	304	18.4	393.39	7.90	21.6	22.02114

H2O: LIME Steps 1 and 2

```
# Step 1: Create an 'explainer' object using training data and model  
explainer = lime::lime(x = as.data.frame(h_train[, features]),  
                      model = model_automl@leader)
```

```
# Step 2: Turn 'explainer' into 'explanations' for test set  
explanations = lime::explain(x = as.data.frame(h_test[, features]),  
                             explainer = explainer,  
                             n_permutations = 5000,  
                             feature_select = "auto",  
                             n_features = 5) # look at top 5 features only
```


H2O: LIME Explanations

```
head(explanations, 5)
```

```
##   model_type case  model_r2 model_intercept model_prediction feature
## 1 regression    1 0.6317351      20.44714      35.20731       b
## 2 regression    1 0.6317351      20.44714      35.20731      rm
## 3 regression    1 0.6317351      20.44714      35.20731     lstat
## 4 regression    1 0.6317351      20.44714      35.20731      dis
## 5 regression    1 0.6317351      20.44714      35.20731      tax
##   feature_value feature_weight  feature_desc
## 1      392.9000     0.001085588 391 < b ≤ 396
## 2       6.8160     9.736177205   6.62 < rm
## 3       3.9500     5.587370832  lstat ≤ 6.95
## 4       8.3248    -1.911507395   5.19 < dis
## 5      256.0000     1.347043279   tax ≤ 279
##
## 1 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 2 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 3 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 4 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
## 5 0.01432, 100.00000, 1.32000, 1.00000, 0.41100, 6.81600, 40.50000, 8.32480, 5.00000, 256.00000, 15.10000, 392.90000
##   prediction
## 1    31.02489
## 2    31.02489
## 3    31.02489
## 4    31.02489
## 5    31.02489
```

H2O: LIME Visualisation

```
# Step 3: Visualise explanations  
lime::plot_features(explanations, ncol = 2)
```



Classification Example

Classification Example: Glass

```
library(mlbench) # for dataset
data("Glass")

# Rename columns
colnames(Glass) = c("Refractive_Index", "Sodium", "Magnesium", "Aluminium",
                    "Silicon", "Potassium", "Calcium", "Barium", "Iron", "Type")
dim(Glass)

## [1] 214 10

str(Glass)

## 'data.frame': 214 obs. of 10 variables:
## $ Refractive_Index: num 1.52 1.52 1.52 1.52 1.52 ...
## $ Sodium          : num 13.6 13.9 13.5 13.2 13.3 ...
## $ Magnesium        : num 4.49 3.6 3.55 3.69 3.62 3.61 3.6 3.61 3.58 3.6 ...
## $ Aluminium        : num 1.1 1.36 1.54 1.29 1.24 1.62 1.14 1.05 1.37 1.36 ...
## $ Silicon          : num 71.8 72.7 73 72.6 73.1 ...
## $ Potassium         : num 0.06 0.48 0.39 0.57 0.55 0.64 0.58 0.57 0.56 0.57 ...
## $ Calcium           : num 8.75 7.83 7.78 8.22 8.07 8.07 8.17 8.24 8.3 8.4 ...
## $ Barium            : num 0 0 0 0 0 0 0 0 0 ...
## $ Iron              : num 0 0 0 0 0 0.26 0 0 0 0.11 ...
## $ Type              : Factor w/ 6 levels "1","2","3","5",..: 1 1 1 1 1 1 1 1 1 1 ...
```

Glass (Simple Split)

```
# Define Features
features = setdiff(colnames(Glass), "Type")
features

## [1] "Refractive_Index" "Sodium"           "Magnesium"
## [4] "Aluminium"        "Silicon"          "Potassium"
## [7] "Calcium"          "Barium"           "Iron"

# Pick four random samples for test dataset
set.seed(1234)
row_test_samp = sample(1:nrow(Glass), 4)
```

H2O AutoML

```
# Start a local H2O cluster (JVM)
library(h2o)
h2o.init(nthreads = -1)

## Connection successful!
##
## R is connected to the H2O cluster:
##   H2O cluster uptime:      3 days 9 hours
##   H2O cluster timezone:    Europe/London
##   H2O data parsing timezone: UTC
##   H2O cluster version:     3.18.0.1
##   H2O cluster version age: 6 days
##   H2O cluster name:        H2O_started_from_R_jofaichow_ydb410
##   H2O cluster total nodes: 1
##   H2O cluster total memory: 3.23 GB
##   H2O cluster total cores: 8
##   H2O cluster allowed cores: 8
##   H2O cluster healthy:     TRUE
##   H2O Connection ip:       localhost
##   H2O Connection port:     54321
##   H2O Connection proxy:    NA
##   H2O Internal Security:  FALSE
##   H2O API Extensions:    XGBoost, Algos, AutoML, Core V3, Core V4
##   R Version:               R version 3.4.3 (2017-11-30)
```

Prepare H2O Data Frames

```
# Prepare Data
h_train = as.h2o(Glass[-row_test_samp,])
h_test = as.h2o(Glass[row_test_samp,])

head(h_test)

##   Refractive_Index Sodium Magnesium Aluminium Silicon Potassium Calcium
## 1          1.51720  13.38      3.50     1.15    72.85      0.50     8.43
## 2          1.51813  13.43      3.98     1.18    72.49      0.58     8.15
## 3          1.52020  13.98      1.35     1.63    71.76      0.39    10.56
## 4          1.52614  13.70      0.00     1.36    71.24      0.19    13.44
##   Barium Iron Type
## 1      0 0.00    1
## 2      0 0.00    2
## 3      0 0.18    2
## 4      0 0.10    2
```

Train Multiple H2O Models

```
# Train multiple H2O models with a simple API
# Stacked Ensembles will be created from those H2O models
# You tell H2O 1) how much time you have and/or 2) how many models do you want
model_automl = h2o.automl(x = features,
                           y = "Type",
                           training_frame = h_train,
                           nfolds = 5,
                           max_runtime_secs = 120, # time
                           max_models = 20,       # max models
                           stopping_metric = "mean_per_class_error",
                           seed = 1234)
```

H2O: AutoML Model Leaderboard

```
# Print out leaderboard
model_automl@leaderboard

##                                     model_id mean_per_class_error
## 1  GBM_grid_0_AutoML_20180219_195906_model_1          0.304868
## 2  GBM_grid_0_AutoML_20180219_195906_model_3          0.304868
## 3  GBM_grid_0_AutoML_20180219_195906_model_2          0.304868
## 4  GBM_grid_0_AutoML_20180219_195906_model_0          0.343727
## 5  GBM_grid_0_AutoML_20180219_195906_model_12         0.347430
## 6 XRT_0_AutoML_20180219_195906                      0.351009
##
## [22 rows x 2 columns]
```

H2O: Model Leader

```
# Best Model (either an individual model or a stacked ensemble)
model_automl@leader

## Model Details:
## =====
##
## H2OMultinomialModel: gbm
## Model ID: GBM_grid_0_AutoML_20180219_195906_model_1
## Model Summary:
##   number_of_trees number_of_internal_trees model_size_in_bytes min_depth
##   1              44                  264          49593           1
##   max_depth mean_depth min_leaves max_leaves mean_leaves
##   1          7      6.10606       2          13      9.99242
##
##
## H2OMultinomialMetrics: gbm
## ** Reported on training data. **
##
## Training Set Metrics:
## =====
##
## Extract training frame with `h2o.getFrame("automl_training_file9053629ae94_sid_920b_12")`  

## MSE: (Extract with `h2o.mse`) 0.01236951  

## RMSE: (Extract with `h2o.rmse`) 0.1112183  

## Logloss: (Extract with `h2o.logloss`) 0.08444432  

## Mean Per-Class Error: 0  

## Confusion Matrix: Extract with `h2o.confusionMatrix(`models`, train = TRUE)`
```

H2O: Making Prediction

```
# Using the best model to make predictions on test set
yhat_test = h2o.predict(model_automl@leader, h_test)
head(yhat_test)

##   predict      p1      p2      p3      p5      p6
## 1      1 0.69766464 0.1366174 0.148579127 0.005158348 0.005166471
## 2      2 0.07421280 0.9013498 0.016233923 0.001808263 0.001817319
## 3      5 0.03841663 0.2580827 0.025503540 0.623656176 0.043140391
## 4      2 0.02248331 0.8882036 0.006976351 0.063091421 0.002743069
##          p7
## 1 0.006813970
## 2 0.004577877
## 3 0.011200592
## 4 0.016502271
```


H2O: LIME Steps 1 and 2

```
# Step 1: Create an 'explainer' object using training data and model
explainer = lime::lime(x = as.data.frame(h_train[, features]),
                      model = model_automl@leader)

# Step 2: Turn 'explainer' into 'explanations' for test set
explanations = lime::explain(x = as.data.frame(h_test[, features]),
                             explainer = explainer,
                             n_permutations = 5000,
                             feature_select = "auto",
                             n_labels = 1, # Explain top prediction only
                             n_features = 5)
```

H2O: LIME Visualisation


```
# Step 3: Visualise explanations  
lime::plot_features(explanations, ncol = 2)
```



Other Stuff

H2O in Action


H₂O Core




H₂O Core




H₂O Core



H₂O Core





Tools & Examples

Python Tools

- lime (Original Python Package by Marco Ribeiro)
[Link](#)

Python Examples

- Marco's Examples [See GitHub README](#)
- LIME + H2O Example [Link](#)
- LIME in Python by Erin Brown [Link](#)

R Examples

- Text Example by Thomas [Link](#)
- HR Analytics Example by Matt [Link](#)
- Cancer Example by Kasia [Link](#)

Related Topics

- SHAP (SHapley Additive exPlanations)
 - A Unified Approach to Interpreting Model Predictions
 - [Paper](#)
 - [GitHub](#)
 - <http://www.f1-predictor.com/model-interpretability-with-shap/>

Amsterdam Meetups

Tue 20 Feb - Sparkling Water in Production Webinar

- Link: <https://www.meetup.com/Amsterdam-Artificial-Intelligence-Deep-Learning/events/247630667/>

Thu 22 Feb - Meetup at ING

- Anomaly Detection in Finance using Isolation Forest by **Andreea Bejinaru**
- FoR the HoRde: WoRld of WaR-and SpArkCRAft by **Vincent Warmerdam**
- Link: <https://www.meetup.com/Amsterdam-Artificial-Intelligence-Deep-Learning/events/247356503/>

Thanks!

joe@h2o.ai / @matlabulous

https://github.com/woobe/lime_water/

Slides created via the R package **xaringan**.